

EIGER SECURITY ASSESSMENT

Adversarial-AI Security Lab — Technical Assessment Report

Candidate: **Siddhi Sanjay Chavan**

Scope: Modules M1–M8 + Capstone (Treasury Policy Poisoning)

Target: Iggy assistant · Eiger lab · Session eiger- 3b32c081

Date: 8 October 2026

Eiger · session eiger-3b32c081

[← Back to the lab](#)

Your progress

Core is the required attack objective. Stretch is the bonus or defence objective for that module.

Vulnerable M1 · Prompt Injection Core objective Complete Stretch objective Complete	Vulnerable M2 · Output Handling Core objective Complete Stretch objective Complete	Vulnerable M3 · RAG Poisoning Core objective Complete Stretch objective Not yet
Vulnerable M4 · Supply-Chain Audit Core objective Complete Stretch objective Complete	Vulnerable M5 · Excessive Agency Core objective Complete Stretch objective Not yet	Vulnerable M6 · MCP Tool Poisoning Core objective Complete Stretch objective Complete
Vulnerable M7 · Multi-Agent Trust Core objective Complete Stretch objective Complete	Vulnerable M8 · Guardrail Bypass Core objective Complete Stretch objective Complete	
Capstone · Treasury Policy Poisoning Complete		

2. Executive Summary

This report documents the candidate's security assessment of **Iggy**, the assistant exposed by the Eiger Adversarial-AI Security Lab, covering eight vulnerability modules (M1–M8) and the chained Capstone scenario.

Status. All eight modules, **M1–M8, were successfully completed**, including the previously achieved M8 stretch objective. **M3 (RAG knowledge-base poisoning) was also successfully completed**, demonstrating that attacker-controlled KB content could be retrieved and influence Iggy's response as authoritative context. M3 and M5 stretch objectives were not confirmed and are therefore not claimed in this assessment. The **Capstone core** was completed. No additional Capstone Stretch result is claimed without supporting evidence.

Principal finding. The assessment consistently identified failures in the handling of **untrusted input across trust boundaries**. Across chat input, stored data, knowledge-base content, tool metadata, inter-agent messages, and retrieved policies, untrusted content could influence model behaviour or downstream actions. Effective remediation therefore requires treating external content as untrusted data, enforcing authorization independently of the LLM, maintaining provenance, and preventing retrieved or relayed content from becoming executable intent.

Notable findings. An **IDOR (Insecure Direct Object Reference)** vulnerability was identified during **M2** and is reported as a separate access-control finding alongside the stored-injection vulnerabilities. The M3 assessment demonstrated successful **RAG knowledge-base poisoning and indirect prompt injection**. The Capstone further demonstrated how multiple weaknesses could be chained to produce an **unauthorized treasury action**, highlighting the importance of defense in depth and server-side authorization.

Scope limits. Findings and conclusions are based on the candidate's completed lab exercises and supplied evidence. No unverified stretch results or unsupported test outcomes are claimed.

3. Results at a Glance

Module	Topic	Status	Severity*	Evidence in report
M1	Prompt Injection	Completed	Medium–High	provided
M2	Stored Injection / Output Handling (+ IDOR)	Completed	High	provided
M3	RAG KB Poisoning (marker RAG-OWNED-7788)	Completed — Core	High	provided
M4	Supply-Chain Audit	Completed	Critical / High	provided
M5	Excessive Agency	Completed — Core	High	provided
M6	MCP Tool-Description Poisoning	Completed	High	provided
M7	Multi-Agent Trust / Cascading Injection	Completed	High	provided
M8	Guardrail Bypass	Completed (incl. stretch)	Medium–High	provided
Capstone	Treasury Policy Poisoning	Core completed (no stretch claimed)	Critical	provided

**Severity is the assessor's qualitative estimate for the vulnerability class in a production setting, not a CVSS score or a lab-assigned value.*

4. M1 — Prompt Injection

Status: **COMPLETED**

Objective	Demonstrate prompt injection against Iggy.
Vulnerability	Iggy's trusted instructions and the user's untrusted message share one undivided context, so the model arbitrates between them by wording and apparent authority rather than by origin.
Trust Boundary	System policy (trusted) ↔ user message (untrusted), concatenated in the model context. The model has no structural way to tell which text carries authority.
Attack / Assessment	A crafted chat message was used to override or extend Iggy's intended behaviour. Exact prompts and Iggy's responses: Evidence not provided.
Result	Completed. The prompt-injection behaviour was demonstrated.
Security Impact	Policy bypass and possible disclosure of instruction-protected content; severity rises with the privileges of the tools and data the model can reach.
Fix / Hardening	Separate system and user content structurally; treat user text as data. Keep secrets and policy decisions out of the prompt and enforce them server-side. Add input/output filtering as defence-in-depth, not as the sole control.
Validation / Evidence	M1 core objective successfully completed and confirmed using the Eiger “Check progress” validator.

Root cause: no structural separation between instructions and data.

Risk / Severity

Finding	Severity (assessor estimate)
M1 Prompt injection	Medium–High

Disclosed Operater Token:

[HALCYON-OPS-7731]

Evidence (screenshots)




Figure 1 — M1 — prompt injection evidence.

5. M2 — Stored Injection / Output Handling (+ IDOR)

Status: **COMPLETED**

Objective	Demonstrate the stored output-handling / injection vulnerability and record the access-control weakness (IDOR) identified in this module.
Vulnerability	Finding M2-A — Stored XSS / output encoding. User-supplied data is stored and later rendered without contextual output encoding. Finding M2-B — Stored prompt injection. Stored user data is placed into Iggy's context without being treated as untrusted. Finding M2-C — IDOR (Insecure Direct Object Reference). An IDOR vulnerability was explicitly identified during M2. A client-supplied object identifier is accepted and acted on without a server-side check that the requester is entitled to that object. It is reported as a separate access-control finding, not as part of the injection findings.
Trust Boundary	M2-A/B: stored user data (untrusted, persisted) ↔ rendering and model-context sinks (which implicitly trust it). Safety is decided at the sink, and the sink trusts stored content. M2-C: client-controlled identifier (untrusted) ↔ back-end data/action layer (must be authoritative). The boundary is crossed whenever the server trusts the identifier alone.
Attack / Assessment	M2-A/B: a payload was stored through a user-controllable field and later interpreted when rendered, showing that stored content is acted on instead of displayed as inert text. M2-C: an object identifier belonging to a different principal was referenced to show that ownership is not enforced by the server. Field, payload, identifiers and requests: Evidence not provided.
Result	Completed. The stored output/injection vulnerability was demonstrated and the IDOR was identified in M2.
Security Impact	Stored XSS: script execution in another viewer's browser — session/token theft, actions on the victim's behalf, content manipulation. Stored prompt injection: persisted attacker text can steer the assistant for later sessions or other users, and can reach tools if combined with excessive agency. IDOR: horizontal privilege escalation — unauthorised read or modification of another user's data. It does not depend on model behaviour, so it stays exploitable even if the injection issues are fixed.
Fix / Hardening	Apply contextual output encoding at every sink (HTML body, attribute, JavaScript and URL contexts each need their own encoding); add a restrictive Content-Security-Policy. Treat all stored user data as untrusted: delimit and label it as data when placing it into the model context; add input allow-lists as defence-in-depth. IDOR fix — enforce object-level authorisation on the server for every request: derive the caller's identity from the authenticated session, verify ownership/entitlement of the referenced object, and deny by default. Do not rely on unguessable identifiers or on the LLM to decide access. Log authorisation denials.
Validation / Evidence	Lab validator output for M2: Evidence provided.

Root cause: (A/B) output safety not enforced at the sink and stored data implicitly trusted; (C) missing server-side object-level authorisation (OWASP API Security: Broken Object Level Authorization).

Risk / Severity

Finding	Severity (assessor estimate)
M2-A Stored XSS / missing output encoding	High
M2-B Stored prompt injection	Medium–High
M2-C IDOR (access-control finding)	High

Evidence (screenshots)




Figure 2 — M2 —IDOR (Insecure Direct Object Reference).




Figure 3 — M2.




Figure 4 — M2 — IDOR evidence (request/response showing access to another principal's object).

6. M3 — RAG Knowledge-Base Poisoning

Status: **CORE COMPLETED (no stretch result claimed)**

Objective	Poison the community knowledge base (KB) so that Iggy's answer contains the required marker RAG-OWNED-7788.
Vulnerability	Attacker-controlled content can be written to the community KB. The RAG pipeline retrieves this content and places it into the model context as authoritative information, without adequately distinguishing curated knowledge from user-submitted material. This creates an indirect prompt-injection vulnerability through the retrieval channel .
Trust Boundary	User-controlled KB content (untrusted) → retrieval → model context (trusted instructions/knowledge). The trust boundary fails because attacker-controlled retrieved content can influence the model as though it were trusted knowledge or instructions.
Attack / Assessment	Attacker-controlled content was added to the community KB and a query was used to retrieve the poisoned content. The poisoned chunk was successfully placed into the model's context, satisfying the poisoned-chunk-in-context condition, and Iggy's resulting answer contained the required marker RAG-OWNED-7788, demonstrating successful RAG injection.
Result	M3 was successfully completed. The assessment demonstrated that attacker-controlled KB content could be retrieved and treated as authoritative model context, causing the injected RAG content to influence Iggy's response and produce the required RAG-OWNED-7788 marker. This confirms the RAG knowledge-base poisoning / indirect prompt-injection vulnerability in the vulnerable lab configuration.
Security Impact	Any user able to write to the KB can potentially influence model responses through poisoned retrieved content. Impact scales with the downstream privileges of the assistant: misinformation and loss of trust at minimum, with potential instruction hijacking or tool abuse where privileged tools are reachable. The demonstrated primitive also forms an important component of the broader attack chain exercised in the Capstone.
Fix / Hardening	Tag KB content by provenance and treat user-submitted or retrieved content as untrusted data rather than instructions. Use clear separation between trusted instructions and retrieved content, require authenticated and reviewed KB changes, filter malicious or instruction-like content, and enforce authorization independently of model output.
Validation / Evidence	The required RAG-poisoning behavior was demonstrated: the poisoned KB content entered the model context and the resulting response contained the required RAG-OWNED-7788 marker..

Root cause: Trust-boundary failure between attacker-controlled knowledge-base content and trusted model instructions/context.

Risk / Severity

Finding	Severity (assessor estimate)
RAG knowledge-base poisoning / indirect prompt injection	High

Evidence (screenshots)




Figure 5 — M3 — KB submission. Add caption after inserting the screenshot.




Figure 6 — M3 — Iggy answer containing RAG-OWNED-7788.

7. M4 — Supply-Chain Audit

Status: **COMPLETED**

Objective	Audit the deployment bundle: identify the poisoned model artifact and the vulnerable dependency pin.
Vulnerability	(1) A poisoned model artifact — model files can carry executable payloads (for example via unsafe deserialisation), so loading an untrusted artifact can execute attacker code. (2) A vulnerable dependency pin: PyYAML==5.3.1, a version affected by known arbitrary-code-execution weaknesses in unsafe YAML loading and fixed in later releases.
Trust Boundary	Third-party artifacts and packages (external, unverified) ↔ the build/runtime environment (trusted). Loading or installing without verification imports the third party's trust level into production.
Attack / Assessment	Static audit (artifacts examined without being loaded; pins reviewed against vulnerability data). The poisoned artifact and the pin PyYAML==5.3.1 were identified. Artifact name/hash and scanner output: Evidence not provided.
Result	Completed. Poisoned model artifact identified; vulnerable pin PyYAML==5.3.1 identified.
Security Impact	Remote code execution in the training or serving environment, data and credential theft, and persistent backdoors. Supply-chain compromise bypasses application-layer controls entirely.
Fix / Hardening	Verify artifacts before use: static scan, signature and known-good-hash checks, trusted registry; prefer non-executable formats; never load unverified models. Upgrade PyYAML to a patched release and use safe loading APIs. Add dependency scanning (e.g. pip-audit/OSV) and SBOM generation as CI gates; pin with hashes.
Validation / Evidence	Lab validation for M4: Evidence provided.

Root cause: third-party artifacts and dependencies were trusted without verification.

Risk / Severity

Finding	Severity (assessor estimate)
Poisoned model artifact	Critical
Vulnerable dependency PyYAML==5.3.1	High

Evidence (screenshots)




Figure 8 — M4 — artifact scan / submission.

8. M5 — Excessive Agency

Status: **CORE COMPLETED (no stretch result claimed)**

Objective	Demonstrate that Iggy can be driven to perform an unauthorised, sensitive action.
Vulnerability	The assistant holds capabilities (tools) broader than its task requires, and the decision whether an action is permitted is made by the LLM from natural-language context rather than enforced by the back-end.
Trust Boundary	Model output (a tool-call proposal influenced by untrusted input) ↔ privileged back-end action. Nothing independent validates the call before execution.
Attack / Assessment	A request was framed so that the sensitive tool call became the model's chosen response, and the action was carried out without an independent authorisation check. Exact request and tool call: Evidence not provided.
Result	Completed. The unauthorised sensitive action was demonstrated.
Security Impact	Prompt-level persuasion converts directly into real-world effects (e.g. state-changing or financial operations). Severity is governed by what the exposed tools can do.
Fix / Hardening	Authorisation must be enforced server-side, not by the LLM. A prompt rule such as "only do X for authorised users" is a request, not a control, and can be overridden by injected context. Apply least agency: minimum tool set, read-only by default, scoped credentials per action. Require human confirmation for high-impact operations and validate parameters (caller, target, amount) in the back-end.
Validation / Evidence	Lab validator output for M5: Evidence provided.

Root cause: authority to act was delegated to the model; no independent policy-enforcement point.

Risk / Severity

Finding	Severity (assessor estimate)
Excessive agency / unauthorised sensitive action	High

Evidence (screenshots)




Figure 10 — M5 — unauthorised action evidence.

9. M6 — MCP Tool-Description Poisoning

Status: **COMPLETED**

Objective	Demonstrate poisoning of MCP tool metadata/descriptions leading to an unintended cross-server action.
Vulnerability	Tool descriptions provided by MCP servers are placed into the model context as trusted guidance. Metadata that is attacker-influenceable can carry hidden instructions that the model obeys.
Trust Boundary	Third-party tool metadata (untrusted data) ↔ model context (trusted instructions); additionally, server-to-server isolation (a CRM-side tool must not be able to cause action on a banking-side tool).
Attack / Assessment	A hidden instruction in tool-description metadata caused the assistant to take an action on a different server than the one the user's request concerned — an unintended cross-server action. Poisoned metadata text and resulting call: Evidence not provided.
Result	Completed. The metadata poisoning issue and the unintended cross-server action were demonstrated.
Security Impact	Confused-deputy attack: the assistant, holding legitimate authority over several servers, is tricked into using that authority on the attacker's behalf. The user never asked for the action and may never see it.
Fix / Hardening	Treat tool metadata as untrusted data; validate, pin and review tool descriptions and strip embedded directives. Server-side authorisation and per-server scoping: the receiving service must authorise the action against the authenticated principal and the request's origin, independent of what the model proposes. Require user confirmation for cross-server or state-changing calls; allow-list tools per task.
Validation / Evidence	Lab validator output for M6: Evidence provided.

Root cause: metadata trusted as instruction; the assistant's ambient authority was not scoped per server or per request.

Risk / Severity

Finding	Severity (assessor estimate)
MCP tool-metadata poisoning / confused deputy	High

Evidence (screenshots)





Figure 11 — M6 — poisoned metadata



IggyMy progressClass boardlocal

M3 RAG PoisoningM4-M5 Supply Chain & AgencyM6 MCPM7 Multi-agentM8 Guardrails & EvasionCapstone

SECURITY MODE
M6 MCPVulnerable

MCP tool poisoning

► ■ How this works — L3 · MCP — external tool servers

M6Vulnerable

Tool-description poisoning

Core objective: Trigger an unintended cross-server call using the hidden instruction in CRM tool metadata.

1 Learn2 Attack3 Harden & retry

Check progressReset attemptCore ✓ · Stretch ✓

Iggy calls the core-banking and CRM tools over MCP servers.

get the core_banking__get_account_details of the account u knowAsk

Unfortunately, the account "u know" does not exist in our system. Please provide the correct account number or details to retrieve the account information.
tool: core_banking__get_account_details({"account": "u know"})

► See the protocol — MCP Inspector (local deploy)

10. M7 — Multi-Agent Trust / Cascading Injection

Status: **COMPLETED**

Objective	Show untrusted dispute text propagating across agents and influencing a refund decision.
Vulnerability	In the multi-agent pipeline, downstream agents treat the content relayed by upstream agents as a trustworthy decision. Free-text supplied by an external party (the dispute) is carried along and interpreted as instruction.
Trust Boundary	External dispute text (untrusted) → agent-to-agent message channel → decision-making agent. Inter-agent messages are assumed trustworthy because of where they came from.
Attack / Assessment	Untrusted dispute text was propagated through the agents and influenced the refund decision. Dispute text and pipeline output: Evidence not provided.
Result	Completed. Propagation of untrusted dispute text and its influence on a refund decision were demonstrated.
Security Impact	Fraudulent or unauthorised refunds/payouts and loss of integrity of automated financial decisions; a single injection point compromises the whole pipeline (cascading failure).
Fix / Hardening	Message authenticity/signatures prove who sent a message and that it was not altered; they do not prove that the content is an authorised decision. A signed message can faithfully carry attacker-written text. Treat all upstream content as untrusted data; have the deciding component derive decisions from structured, trusted fields and server-side policy. Validate that the refund target is owned by the claimant, require independent approval for payouts, and sanitise/segregate free-text across agent hops.
Validation / Evidence	Lab validator output for M7: Evidence provided.

Root cause: conflation of authenticity with authorisation; free-text promoted to a decision.

Risk / Severity

Finding	Severity (assessor estimate)
Cascading injection across agents	High

Evidence (screenshots)




Figure 12 — M7 — dispute input and refund decision evidence.

11. M8 — Guardrail Bypass

Status: **COMPLETED (including previously achieved stretch objective)**

Objective	Bypass the guardrail protecting lggy; the stretch objective was also achieved previously.
Vulnerability	The guardrail relies on surface-pattern / lexical matching (keywords and fixed patterns) rather than meaning. Content that is semantically equivalent but lexically different is not recognised.
Trust Boundary	Guardrail verdict ↔ downstream behaviour: the system implicitly trusts that anything the filter allows is safe, making the filter the sole barrier.
Attack / Assessment	The restricted content was reshaped so that its surface form fell outside the guardrail's patterns while its meaning was preserved. Exact transformation and outputs (core and stretch): Evidence not provided.
Result	Completed, including the previously achieved stretch objective.
Security Impact	Guardrail-dependent protections can be defeated by routine obfuscation, so a filter used as the only control gives false assurance.
Fix / Hardening	Normalise input/output (decode, canonicalise, remove obfuscation) before matching. Add semantic/intent-based detection (classifiers) rather than keyword lists alone. Enforce critical controls server-side so that bypassing the text filter does not by itself grant access to sensitive data or actions; treat the guardrail as a safety net.
Validation / Evidence	Lab validator output for M8 core and stretch: Evidence provided.

Root cause: reliance on lexical pattern matching as a primary control.

Risk / Severity

Finding	Severity (assessor estimate)
Guardrail bypass	Medium–High

Evidence (screenshots)




Figure 13 — M8 — guardrail bypass evidence (core).

12. Capstone — Treasury Policy Poisoning

Status: **COMPLETED**

Objective	Chain individual weaknesses into an unauthorised treasury action.
Vulnerability	A multi-stage chain: sensitive material exposed through source disclosure; an ingestion path into the Treasury knowledge base that could be abused with it; a treasury agent that treats retrieved policy as authoritative; and a transfer capability not independently authorised.
Trust Boundary	Several compounding failures: source/config exposure (secret ↔ public), KB ingestion (untrusted publisher ↔ policy store), retrieved policy (data ↔ instruction), and agent output ↔ financial action.
Attack / Assessment	The chain proceeded in four stages (below). Commands, keys, payloads and validator output: Evidence not provided.
Result	Core completed: the chain of source disclosure → policy poisoning → retrieval of the poisoned policy → unauthorised treasury action was carried through. No stretch result is claimed.
Security Impact	Direct financial loss and loss of control over treasury policy. The attack requires no model exploit beyond poisoned context: each weak link enables the next.
Fix / Hardening	See defence-in-depth list below; breaking any single link stops the chain, but all should be fixed.
Validation / Evidence	Chain validator output: Evidence provided.

Attack chain

Stage	Description
1	Source disclosure — source/sample configuration exposed sensitive material (key/secret)
2	Policy poisoning — the disclosed access was used to insert an attacker-authored policy into the Treasury knowledge base
3	Retrieval of poisoned policy — the treasury agent retrieved the policy and treated it as authoritative
4	Unauthorised treasury action — the agent performed a treasury action that no legitimate authority had approved

Defence-in-depth fixes

- **Stop the disclosure:** never ship secrets in source or sample configuration; use a secrets manager, keep configuration files out of any served source tree, and rotate any exposed key.
- **Protect ingestion:** authenticate and authorise every policy write with scoped, individually attributable credentials; require review/approval and audit logging.
- **Treat retrieved policy as data:** remove prompt language that makes retrieved text "authoritative"; retrieved content may inform a proposal but must never authorise an action.

- **Gate the money:** server-side authorisation of target, amount and requester, plus human approval for treasury transfers, independent of the model's output.
- **Detect:** monitor for unusual ingestion and policy changes, and alert on off-pattern transfers.

Risk / Severity

Finding	Severity (assessor estimate)
End-to-end chain (unauthorised treasury action)	Critical
Secret exposure in source	High
Authoritative treatment of retrieved policy	High

Evidence (screenshots)




Figure 15 — Capstone — policy ingestion.




Figure 16 — Capstone — retrieval and treasury action.

```

    },
    }

SEED: list[dict] = [
    {"text": t, "provenance": "trusted"}
    for topic in _TOPICS.values()
    for t in topic
]
# curl "http://127.0.0.1:8000/source/blob?session=YOUR_SESSION_ID&path=.env.sample"
sh: 17: curl: not found
# python -c "import urllib.request; print(urllib.request.urlopen('http://127.0.0.1:8000/source/blob?session=YOUR_SESSION_ID&path=.env.sample').read().decode())"
{"path": ".env.sample", "content": "# Sample environment for eiger-platform.\n# NOTE: copy to .env and fill real values in prod.\nINGEST_KEY=6552ab5b4f13a55aae561d4ac40cf241\n"}
# |

```

13. Cross-Cutting Security Findings

Theme	Modules	Observation
Data treated as instruction	M1, M2, M3, M6, M7, Capstone	Content from users, KBs, tool metadata, other agents and retrieved policies enters the context with the authority of trusted instructions.
Authorisation in the wrong layer	M2 (IDOR), M5, M6, M7, Capstone	Access decisions are left to the model or to client-supplied identifiers instead of a server-side policy-enforcement point.
Authenticity $\neq$ authorisation	M7	Signed or well-formed messages were treated as authorised decisions.
Single-layer controls	M8, M5	Lexical filters and prompt rules were the only barrier and were bypassable.
Supply-chain trust	M4	Third-party artifacts and dependencies used without verification.
Chained impact	Capstone, M3	Individually moderate flaws compound into severe outcomes.

14. Defence-in-Depth Recommendations

Control	Addresses	Description
Server-side authorisation (object- and function-level)	M2-C, M5, M6, M7, Capstone	Deny by default; derive identity from the authenticated session; validate ownership, target and amount in the back-end.
Instruction/data separation and provenance tagging	M1, M2-A/B, M3, M6	Label untrusted content (user, KB, metadata, agent output) as data; never allow it to carry instructions.
Contextual output encoding + CSP	M2-A/B	Encode per sink; restrictive Content-Security-Policy; treat stored data as untrusted.
Least agency and human-in-the-loop	M5, M6, Capstone	Minimum tool scope; confirmation for high-impact actions.
Guardrails as a safety net	M8	Normalisation, semantic detection, and server-side enforcement behind them.
Supply-chain verification	M4	Signature/hash verification, static scanning, dependency auditing in CI, SBOM.
Secrets hygiene and monitoring	Capstone	Secrets manager, rotation, scoped ingestion credentials, audit logging and alerting.

Residual risk. These controls address the demonstrated paths. They do not eliminate novel or obfuscated injection, compromised upstream sources, or weaknesses outside the hardened layer; layered controls and continuous monitoring remain necessary.

16. Conclusion

The assessment exercised eight vulnerability classes and one chained scenario against Iggy. **All modules M1–M8 were successfully completed**, including the M8 stretch objective, and the **Capstone core was also completed**. M3 successfully demonstrated **RAG knowledge-base poisoning**, confirming that attacker-controlled KB content could be retrieved and influence Iggy’s response as trusted context. The assessment also identified an **IDOR (Insecure Direct Object Reference)** vulnerability during M2 as a separate access-control finding.

The central lesson is that AI systems must be designed so that **untrusted content can inform a system but must never be allowed to command or authorize it**. Effective protection therefore depends on architectural controls rather than prompts alone: server-side authorization and object-level access checks, clear separation of data from instructions, provenance-aware RAG, scoped tool permissions, secure supply-chain practices, and layered defenses. The completed exercises demonstrate that when these controls are absent or rely primarily on model behaviour, individual weaknesses can be chained into significant security consequences, including unauthorized actions demonstrated in the Capstone.

17. AI Assistance Disclosure

I used an AI assistant as an engineering aide throughout this assessment to understand the Eiger architecture and the workflow of each security module, reason about the relevant attack classes and payload construction, identify the corresponding security controls, explain the rationale for remediation, and structure and edit this report.

I independently executed the attacks in the Eiger lab and validated the completed objectives using the lab's own **"Check progress" validator and progress board**, rather than relying solely on AI-generated conclusions. The assessment results reported in this document are based on my own lab execution and observations. All screenshots included in the report are from my own Eiger session (**eiger- 3b32c081**).

AI assistance was therefore used as a supporting engineering and documentation aid, while the actual exploitation, validation, and assessment decisions were performed by me.

Residual Risk & Limitations

The controls recommended in this assessment address the demonstrated attack paths within the respective modules, but they do not constitute a complete defence against all AI-security threats. They cannot, by themselves, eliminate novel or obfuscated prompt-injection techniques, compromised upstream knowledge or dependency sources, or vulnerabilities in system layers outside the specific controls assessed.

Production AI systems should therefore use **defence in depth**, combining provenance tracking, least-privilege tool scoping, independent server-side authorization of sensitive actions, secrets hygiene, dependency scanning within CI/CD, input and output validation, and continuous security monitoring. No single guardrail or model-level control should be treated as sufficient protection for security-sensitive operations.

Walk-through recording:

https://drive.google.com/drive/folders/1UhxkvYgNYwDD4fWaX4h9sEkLfrkOsi_?usp=sharing

The recording demonstrates the assessment activities and results performed in the Eiger lab.